

---

# ADMINISTRATOR'S GUIDE: BMI USER ACCEPTANCE TESTING FRAMEWORK

---

Massachusetts Open Cloud

Paul Grosu (*pgrosu@gmail.com*)

FIRST EDITION  
2017  
(Revised on 10/10/2017)

# CONTENTS

---

CONTENTS    ii

<b>I</b>	<b>AN APPROACH TO DELIVERING DEFECT-FREE SOFTWARE PRODUCTS</b>	<b>1</b>
1	DOCUMENT-DRIVEN TESTING FOR OPTIMAL DEPLOYMENTS	3
1.1	How to guarantee a defect-free software product	3
1.2	Testing Under Changing Environments Via System Testing	4
1.3	Behavior-Driven Development: A Scientific-Method Approach to Testing	5
<b>II</b>	<b>GETTING STARTED WITH THE USER ACCEPTANCE TESTING FRAMEWORK</b>	<b>9</b>
2	THE USER ACCEPTANCE TESTING FRAMEWORK	11
2.1	The User Acceptance Testing Architecture	11
2.1.1	Configuring a Testing Environment	12
2.1.2	Structure of Test Directory	13
2.1.3	The config file	13
2.2	Preparing The Environment	14
2.3	Performing the Acceptance Tests	15
2.3.1	Entering an Interactive Session	17
INDEX		19

# I

## AN APPROACH TO DELIVERING DEFECT-FREE SOFTWARE PRODUCTS



# DOCUMENT-DRIVEN TESTING FOR OPTIMAL DEPLOYMENTS

---

*The fundamental principle of science, the definition almost, is this:  
the sole test of the validity of any idea is experiment.  
— Richard P. Feynman*

## 1.1 HOW TO GUARANTEE A DEFECT-FREE SOFTWARE PRODUCT

A SOFTWARE PRODUCT CAN BE DEEMED DEFECT-FREE if all the requirements have been completely covered by tests supporting the specifications, and are continuously validated throughout the development cycle. A development team always strives to provide such guarantees, which can be achieved by being *diligent* in following specific software development practices and ensuring that the requirements are bounded and reflected through *complete coverage*.

One methodology of ensuring that such a continuity is preserved, is via the V-Model of software development<sup>1</sup> illustrated in Figure 1.1.

---

<sup>1</sup> <http://ccd.docs.berkeley.edu/content/system-validation-plan>

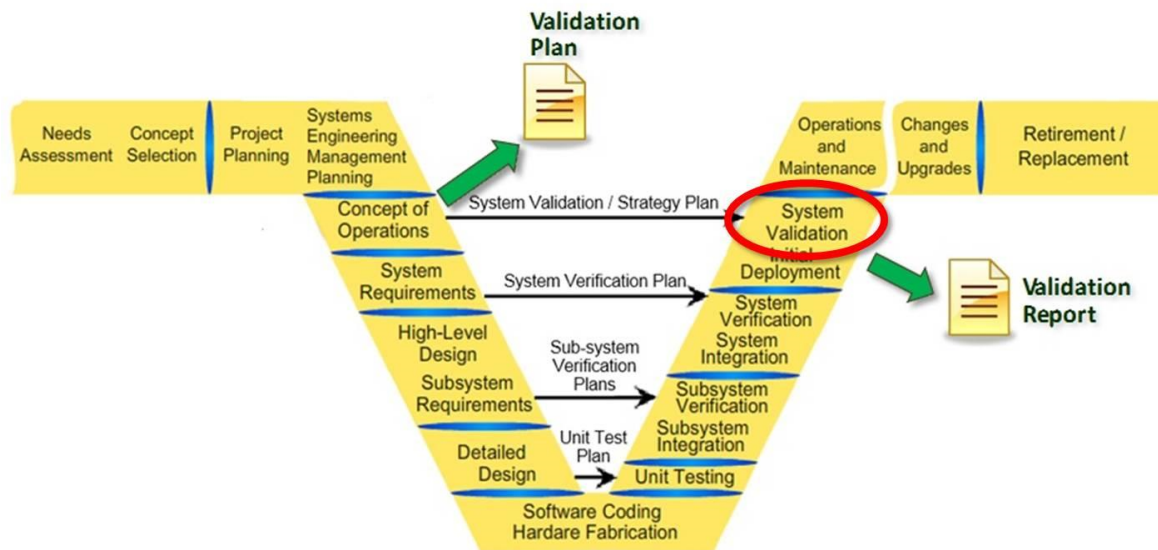


FIGURE 1.1 – The V-Model of software development.

If you bisect vertically the V-Model, you will notice that each requirement – or implementation – on the left-hand side is *supported* by a verification or validation step on the right-hand side. This ensures that the scope is bounded and provides complete coverage<sup>2</sup>. For every step documentation is critical to both *specify* the requirements, and then to subsequently *validate* against those requirements.

## 1.2 TESTING UNDER CHANGING ENVIRONMENTS VIA SYSTEM TESTING

There will be times where *Integration Testing* is not enough. This is where *System Testing*<sup>3</sup> comes in. Here we take the software product as a *black box* – as opposed to in *Integration Testing* – and test it under different environments without touching the code-base. One of the best ways to ensure that the software product will operate as defined by the *requirements* is to run *end-to-end scenarios* with validation. This requires one to have a list of *functional specifications* that the software product must perform, and to create one or more workflows where these will be pipelined together to generate this type of *Functional Testing*<sup>4</sup>.

For BMI these are defined as follows:

- pro ► *Provisions a node.*
- dpro ► *Devisions a node.*

<sup>2</sup>Coverage ensures all aspects of the codebase are verified and asserted via test(s).

<sup>3</sup>Ashfaqe Ahmed and Bhanu Prasad. 2016. *Foundations of Software Engineering*. Auerbach Publications, Boston, MA, USA.

<sup>4</sup>Functional Testing validates the software design based on the requirement specifications, by running tests to check that the software's features match the functional specifications.

- snap ► Takes a snapshot of a node.
- ls ► Lists store images.
- import ► Importing images or snapshots into BMI for provisioning.
- db ► Database commands that about imported images or snapshots.

By then integrating these into an end-to-end workflow, one can perform all these and ensure that the basic requirements are satisfied. An example of a possible end-to-end workflow is described in Figure 1.2.

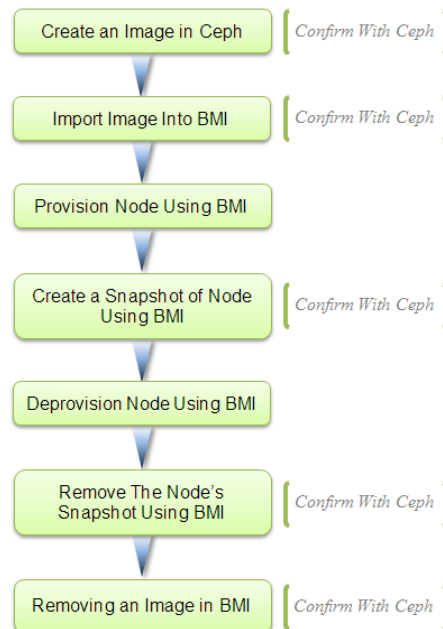


FIGURE 1.2 – An example of an end-to-end workflow.

### 1.3 BEHAVIOR-DRIVEN DEVELOPMENT: A SCIENTIFIC-METHOD APPROACH TO TESTING

The *Scientific Method* is an unbiased approach to discovering what the facts truly about a system by progressing using *systematic doubt*<sup>5</sup> to ensure adequate evidence support each problem being solved. The facts are not gathered unless there is a problem being defined upon which relevant facts are required to prove or disprove inquiries (*hypotheses*) about the problem.

<sup>5</sup>Morris Raphael Cohen and Ernest Nagel. 1934. *An Introduction to Logic And Scientific Method*. Harcourt, Brace and World, New York, NY, USA.

Behavior-Driven Development (BDD) is defined through a live document implemented using the Gherkin language<sup>6</sup>, which utilizes *Given-When-Then* control-flow syntax defined as follows:

- Given ▶ *Defines a given state.*
- When ▶ *Defines a given action performed under the given state.*
- Then ▶ *Defines the expected outcome after the action is performed.*

By building a *scenario* through combining these into premises using the *Given-* and *When-*initiated statements, we are able to discover if our system is validated at each step and confirm the *Then* conclusion statement(s). Thus we are hypothesis-driven through a BDD-model of our system to ensure it matches our expected *operational semantics*<sup>7</sup>.

An example of such an end-to-end scenario for BMI is illustrated in Figure 1.3, where each line is a step that references an implemented function.

The end-to-end scenario is a *model* used as a set of *rules of inference* guided by an ordered collection of *premises* – which are assumed to be *true* – and conclude that *all* the steps are *truth preserving*:

*Premises*  $\implies$  *Conclusions*

---

<sup>6</sup> <https://github.com/cucumber/cucumber/wiki/Gherkin>

<sup>7</sup> [https://en.wikipedia.org/wiki/Operational\\_semantics](https://en.wikipedia.org/wiki/Operational_semantics)



```

Feature: Running an end-to-end acceptance test

Scenario: Importing/Removing Image, DB/Ceph consistency

Given RBD will create an image
  | image_name |
  | bmi-test-image |

And BMI log line-count will be measured at the beginning

When BMI will import an image
  | image_name | project_name |
  | bmi-test-image | bmi_infra |

And BMI will provision a node
  | image_name | project_name | network_name | node_name | NIC |
  | bmi-test-image | bmi_infra | bmi-provision-dev | cisco-05 | enp130s0f0 |

Then BMI will create a snapshot of a node
  | project_name | node_name | snapshot_name |
  | bmi_infra | cisco-05 | bmi-test-image-snapshot |

Then RBD will confirm the snapshot exists
  | snapshot_name |
  | bmi-test-image-snapshot |

And BMI will remove a snapshot
  | snapshot_name | project_name |
  | bmi-test-image-snapshot | bmi_infra |

Then BMI will deprovision a node
  | project_name | network_name | node_name | NIC |
  | bmi_infra | bmi-provision-dev | cisco-05 | enp130s0f0 |

And BMI will remove an image
  | image_name | project_name |
  | bmi-test-image | bmi_infra |

Then RBD will confirm the removed image's clone
  | image_name | project_name |
  | bmi-test-image | bmi_infra |

And RBD will remove the created image
  | image_name |
  | bmi-test-image |

And BMI log line-count will be measured at the end

```

FIGURE 1.3 – The BMI End-to-End Behavior-Driven Deployment Test, with tables of parameters to test with.

For example, in Figure 1.4 the creation of a RADOS block device (RBD<sup>8</sup>) mountable image at the start is defined through the `rbdc_create_image()` function, where it is decorated by the sentence referenced in the live-document.

<sup>8</sup>Ceph Storage provides the ability for its (bootable) images to be mountable remotely using a RADOS block device. For more information please proceed to the following web location:

<https://docs.openstack.org/mitaka/config-reference/block-storage/drivers/ceph-rbd-volume-driver.html>

```

import behave
import time # Needed for Ceph Hammer client consistency
from bmi_config import RBD_CREATE, IMAGE_NAME, PROVISIONING_DELAY
from subprocess import check_output, CalledProcessError, STDOUT
from test_operation import test_event_store_insert, test_rollback

@step('RBD will create an image')
def rbd_create_image(context):
    for row in context.table:
        try:
            print( "      -> Checking that no pre-existing " +
                  row['image_name'] + " is present in Ceph, before creating it...")
            rbd_filename_check_stdout = check_output('rbd ls | grep ' + row['image_name'],
                                                    stderr=STDOUT, shell=True)

        except CalledProcessError:
            pass # The image already exists, as it was previously created

        try:
            print( "      -> Creating the " + row['image_name'] + " image in Ceph...")
            rbd_create_stdout = check_output( 'rbd create ' +
                                             row['image_name'] +
                                             ' --size 1 --image-format 2', stderr=STDOUT, shell=True)

        except CalledProcessError:
            pass # The image already exists, as it was previously created
        except Exception:
            test_rollback(context)
            print( "      -> Checking that " + row['image_name'] + " exists in Ceph...")
            rbd_filename_check_stdout = check_output( 'rbd ls | grep ' +
                                                    row['image_name'], stderr=STDOUT, shell=True)
            context.rbd_filename_check = rbd_filename_check_stdout.strip()

        # Journal the event for rollback
        test_event_store_insert(context, { RBD_CREATE: {IMAGE_NAME : 'image_name'} })

    time.sleep( PROVISIONING_DELAY ) # Needed for Ceph Hammer client consistency
    assert context.rbd_filename_check == row['image_name']

```

FIGURE 1.4 – The definition of the RBD creation step, where the decoration highlights the sentence referenced in the end-to-end deployment test.

In the next chapter, you will learn how to configure and run an acceptance test.

# II

## GETTING STARTED WITH THE USER ACCEPTANCE TESTING FRAMEWORK



# THE USER ACCEPTANCE TESTING FRAMEWORK

---

*This chapter will guide through the steps of creating and running a BDD scenario for BMI via the User Acceptance Framework.*

## 2.1 THE USER ACCEPTANCE TESTING ARCHITECTURE

You will need to be provided a compressed (*tar.gz*) file of the acceptance tests. After you uncompress it via the "`tar -xzf acceptance-tests.tar.gz`" you will see the following files and folders in the root directory:

Name	Date modified	Type	Size
bdd	9/12/2017 9:16 PM	File folder	
config	9/12/2017 9:16 PM	File folder	
doc	9/13/2017 12:24 PM	File folder	
scripts	9/12/2017 9:16 PM	File folder	
test-results	9/12/2017 9:16 PM	File folder	
bmi-uat.py	9/11/2017 2:06 AM	PY File	3 KB
interactive-session_SOURCE-THIS.sh	7/14/2017 2:32 PM	SH File	1 KB
prepare-environment.sh	9/11/2017 4:54 PM	SH File	1 KB
run-multiple-rounds_stress-test.sh	8/22/2017 11:11 PM	SH File	1 KB

FIGURE 2.1 – Root directory of the User Acceptance Testing framework.

Below are descriptions of the critical folders and files required for configuring and for running the tests:

- `config` ▶ Contains the testing configurations for different environments (i.e. PRB, NEU, etc).
- `doc` ▶ Contains this manual.
- `scripts` ▶ Contains the configuration scripts that are run for different stages during testing.
- `prepare-environment.sh` ▶ Configurations to run for different OS environments before testing, which can be used to create cleanup scripts.
- `bmi-uat.py` ▶ The command-line interface (CLI) for listing and running the tests.
- `test-results` ▶ Provides test results in case one performs randomized tests for multiple rounds.

The `interactive-session_SOURCE-THIS.sh` script is used if you want to drop into an interactive session into the environment of a particular test after is completed in order to inspect or rollback changes.

Next you will learn about the `config` directory regarding how to use or create new testing environments.

### 2.1.1 Configuring a Testing Environment

If you look at the `config` directory you will see something that looks similar to this:

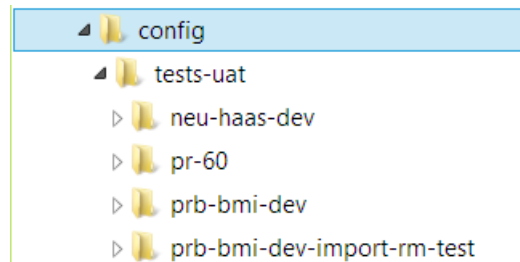


FIGURE 2.2 – `Config` directory of the User Acceptance Testing framework.

It is best to copy a previous directory of interest if you would like to perform the minimal changes to a test. There are two types of test directories:

- Pull-Request Test** ▶ *Performs tests on a specific pull-request (i.e. pr-60). These are usually performed in preparation for running a deployment test.*
- Repository Test** ▶ *Performs tests on a whole repository (i.e. neu-haas-dev). These would be performed to ensure a release is ready for deployment.*

Next we will look at how a test configuration is structured.

### 2.1.2 Structure of Test Directory

If you look at any of the test directories they all look as follows:

Name	Date modified	Type	Size
doc	9/12/2017 9:16 PM	File folder	
features	9/12/2017 9:16 PM	File folder	
scripts	9/12/2017 9:16 PM	File folder	
steps	9/12/2017 9:16 PM	File folder	
bmi-config.sh	7/14/2017 6:41 PM	SH File	1 KB
config	9/11/2017 2:04 AM	File	1 KB
customize	7/14/2017 10:09 PM	File	1 KB
customize-after-git-clone	7/14/2017 5:57 PM	File	1 KB

FIGURE 2.3 – The test directory structure.

To keep the configurations simple and practical, it is important to know about the following four components:

- `config` ► *This file configures the BMI UAT test for the environment, and is the most important file.*
- `bmi-config.sh` ► *This is the second most important file, and is used to configure the BMI pre-test deployment directories.*
- `features` ► *These contain the live documents that can be changed with the exception of the template file. Additional scenario files to test for can be added if preferred.*
- `steps` ► *Contains the functions that map to the given BDD definition in the feature files that build up the scenarios.*

### 2.1.3 The `config` file

The `config` file is usually the only file one will usually configure the most of the time, and it was created to ensure minimal changes are necessary for test-preparation. The structure of the file is shown in Figure 2.4, and is composed of the following three main sections:

- `BMI_RELEASE_NAME` ► *This will denote the name of the directory for the scenario that is being tested, underneath which the tests will be installed, configured and run.*

E2E Test Configs ► *The middle section contains the End-To-End configuration information that are pertinent to the environment being tested (i.e. HIL project names, names of BMI images to create, etc).*

BMI and HIL Configs ► *These contain the HIL and BMI local configurations in order for the tests to run.*

```
export BMI_RELEASE_NAME=moc-0.5-release

export BMI_PROJECT=bmi_infra
export HIL_NODE=cisco-05
export HIL_NIC=enp130s0f0
export HIL_NETWORK=bmi-provision-dev
export BMI_IMAGE_NAME=bmi-test-image
export BMI_SNAPSHOT_NAME=bmi-test-image-snapshot

export BMI_CONFIG=/etc/bmi/bmiconfig_pgrosu.cfg
export HIL_ENDPOINT=http://127.0.0.1:8000
export HIL_USERNAME=*****
export HIL_PASSWORD=*****
```

FIGURE 2.4 – The test configuration file structure.

An example of the `bmi-config.sh` file is shown in Figure 2.4, which provides the configurations of where the BMI instance will be installed (`BMI_INSTANCE_DIR`), and the location of the User Acceptance Tests directory (`ACCEPTANCE_TESTS_SRC_DIR`).

```
export BMI_INSTANCE_DIR=${HOME}/pgrosu/ims-instance
export ACCEPTANCE_TESTS_SRC_DIR=${HOME}/pgrosu/acceptance-tests
export BDD_DIR=./bdd
export BDD_STEPS_DIR=${BDD_DIR}/steps
```

FIGURE 2.5 – An example of a `bmi-config.sh` file.

## 2.2 PREPARING THE ENVIRONMENT

Sometimes the operating environment requires extra functionality – such as Python or Git availability – to be available before running a test. These configurations can be placed as Bash scripts under the `scripts\prepare-environments` directory, as shown in Figure 2.6.





Name	Date modified	Type	Size
 prepare-centos-with-git.sh	8/21/2017 5:13 PM	SH File	1 KB
 prepare-ubuntu-with-python.sh	8/21/2017 4:49 PM	SH File	1 KB

FIGURE 2.6 – An example of the prepare-environments directory.

To list all configurations, type under the main acceptance-tests directory `./prepare-environment.sh`, as shown in Figure 2.9.

```
ubuntu@pgrosu-ubuntu16:~/acceptance-tests$ ./prepare-environment.sh
Please choose one of the following configurations based on your environment:
prepare-centos-with-git
prepare-ubuntu-with-python
ubuntu@pgrosu-ubuntu16:~/acceptance-tests$
```

FIGURE 2.7 – Listing the prepare-environments configurations.

To run a configuration just use the following format to run the appropriate configuration for your environment:

```
./prepare-environment.sh CONFIGURATION
```

Now you are ready to run a test configuration.

### 2.3 PERFORMING THE ACCEPTANCE TESTS

The performance tests can be initiated via the following steps:

1. To list the testable BMI service configurations, type the following command:

```
./bmi-uat.py ls
```

You should see something like the following:

The available configurations are:

```
neu-haas-dev
```

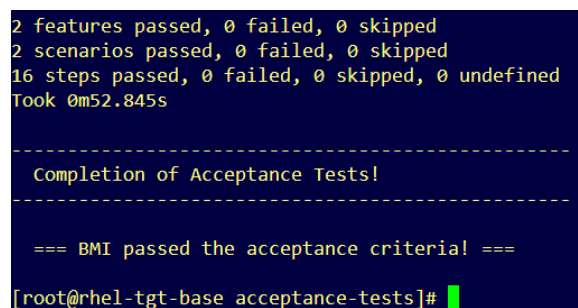
2. To run the standard end-to-end configuration, type the following command:

```
./bmi-uat.py --run BMI_SERVICE_CONFIGURATION
```

Example:

```
./bmi-uat.py --run neu-haas-dev
```

At the end you if the tests passed successfully, you should see the following output:



```
2 features passed, 0 failed, 0 skipped
2 scenarios passed, 0 failed, 0 skipped
16 steps passed, 0 failed, 0 skipped, 0 undefined
Took 0m52.845s

-----
Completion of Acceptance Tests!
-----

=== BMI passed the acceptance criteria! ===

[root@rhel-tgt-base acceptance-tests]#
```

FIGURE 2.8 – The BMI completed successfully the end-to-end scenario.

3. To run the tests with randomized parameters, type the following where the value indicates the number of times to run the test:

```
./bmi-uat.py --run neu-haas-dev --randomize 3
```

4. To check if the tests passed or failed, type the following:

```
./bmi-uat.py check
```

You should see the following:

```
All tests passed!
```

This command checks the `test-results` directory for any subdirectory containing `FAIL` in its name.

5. To cleanup all previous results, type the following:

```
./bmi-uat.py clean
```

You will notice that when running a test, there are many additional sanity-checks that are being made to ensure each test not only completes properly, but also provides sufficient detail in case of failure, as shown by the following figure:

```

When BMI will import an image                                     # tests/bdd/steps/bmi_import.py:7
| image_name | project_name |
| bmi-test-image | bmi_infra |
Running: bmi import bmi_infra bmi-test-image
-> Checking that bmi-test-image exists in Ceph...
Running: bmi db ls | grep bmi-test-image
-> Checking that 112233445566778899img2 exists in BMI's database...
Running: rbd ls | grep 112233445566778899img2

```

FIGURE 2.9 – Sanity-checks for the BMI import step.

### 2.3.1 Entering an Interactive Session

After the test has completed – either successfully or not – one can enter an interactive session to inspect the state of the test, where BMI commands can be executed interactively. This is performed from the acceptance-tests directory by typing the following command – make sure to not forget the period (.) at the before the script-name:

```
. interactive-session_SOURCE-THIS.sh
```

If the test was based on a pull-request, please add `-pr` as follows:

```
. interactive-session_SOURCE-THIS.sh -pr
```

You should see something as follows, showing that one is in a virtual environment:

```
(.bmi_venv) ubuntu@pgrosu-ubuntu16:~/ims-instance/ims$
```

To exit the interactive session, just type the following command – from within the session:

```
. return-back-to-acceptance-dir_SOURCE-THIS.sh
```

By following the above steps you can now test your own customizations of BMI any services.



# INDEX

---

BMI and HIL Configurations, 14  
BMI\_RELEASE\_NAME, 13  
E2E Test Configs, 14  
Pull-Request Test, 12  
Repository Test, 12  
bmi-config.sh script, 13

BDD (Given), 6  
BDD (Then), 6  
BDD (When), 6  
BDD features, 13  
BDD steps, 13  
Behavior-Driven Development (BDD), 5  
bmi-uat.py script, 12

complete coverage, 3  
conclusions, 6  
config directory, 12, 13

doc directory, 12  
end-to-end scenario, 6  
functional testing, 4  
interactive session, 17  
premises, 6  
prepare-environment.sh script, 12  
running User Acceptance Tests, 15  
scientific method, 5  
scripts directory, 12  
system testing, 4

test-results, 12  
V-Model, 3